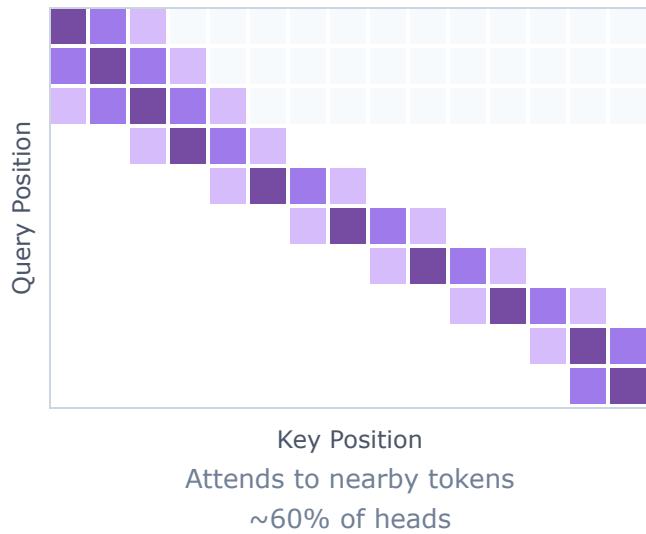
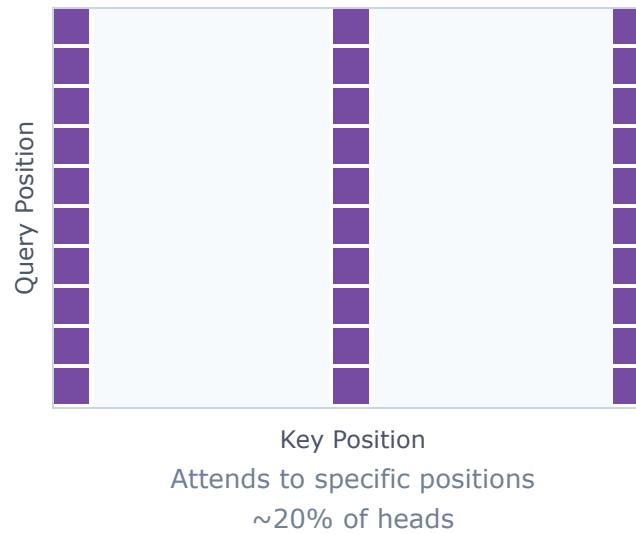


Learned Attention Patterns

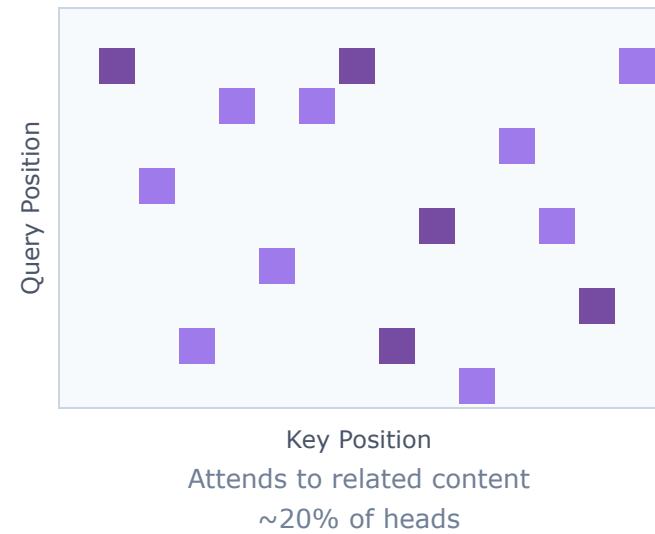
Local Attention



Positional Attention



Semantic Attention



Attention Weight



Optimization Opportunity

Local patterns enable sparse attention: skip low-weight computations
50-80% computation reduction with <1% accuracy impact